

Name :- Ansh Agrawal

Roll NO : 20

Practical 5

Import Libraries

```
In [20]: import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error, r2_score
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
```

Collect datasets

```
In [22]: df = pd.read_csv("banking.csv");
```

```
In [23]: df.head()
```

```
Out[23]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week
0	44	blue-collar	married	basic.4y	unknown	yes	no	cellular	aug	thu
1	53	technician	married	unknown	no	no	no	cellular	nov	fri
2	28	management	single	university.degree	no	yes	no	cellular	jun	thu
3	39	services	married	high.school	no	no	no	cellular	apr	fri
4	55	retired	married	basic.4y	no	yes	no	cellular	aug	fri

5 rows × 21 columns

```
In [24]: df.tail()
```

Out[24]:

	age	job	marital	education	default	housing	loan	contact	month	day_o
41183	59	retired	married	high.school	unknown	no	yes	telephone	jun	
41184	31	housemaid	married	basic.4y	unknown	no	no	telephone	may	
41185	42	admin.	single	university.degree	unknown	yes	yes	telephone	may	
41186	48	technician	married	professional.course	no	no	yes	telephone	oct	
41187	25	student	single	high.school	no	no	no	telephone	may	

5 rows × 21 columns

In [25]:

df.describe()

Out[25]:

	age	duration	campaign	pdays	previous	emp_var_rate	cons_price
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767

In [26]:

df.shape

Out[26]:

(41188, 21)

In [27]:

df.dtypes

```
Out[27]: age                int64
         job                object
         marital            object
         education          object
         default            object
         housing            object
         loan               object
         contact            object
         month              object
         day_of_week        object
         duration           int64
         campaign           int64
         pdays             int64
         previous           int64
         poutcome           object
         emp_var_rate       float64
         cons_price_idx     float64
         cons_conf_idx      float64
         euribor3m          float64
         nr_employed        float64
         y                  int64
         dtype: object
```

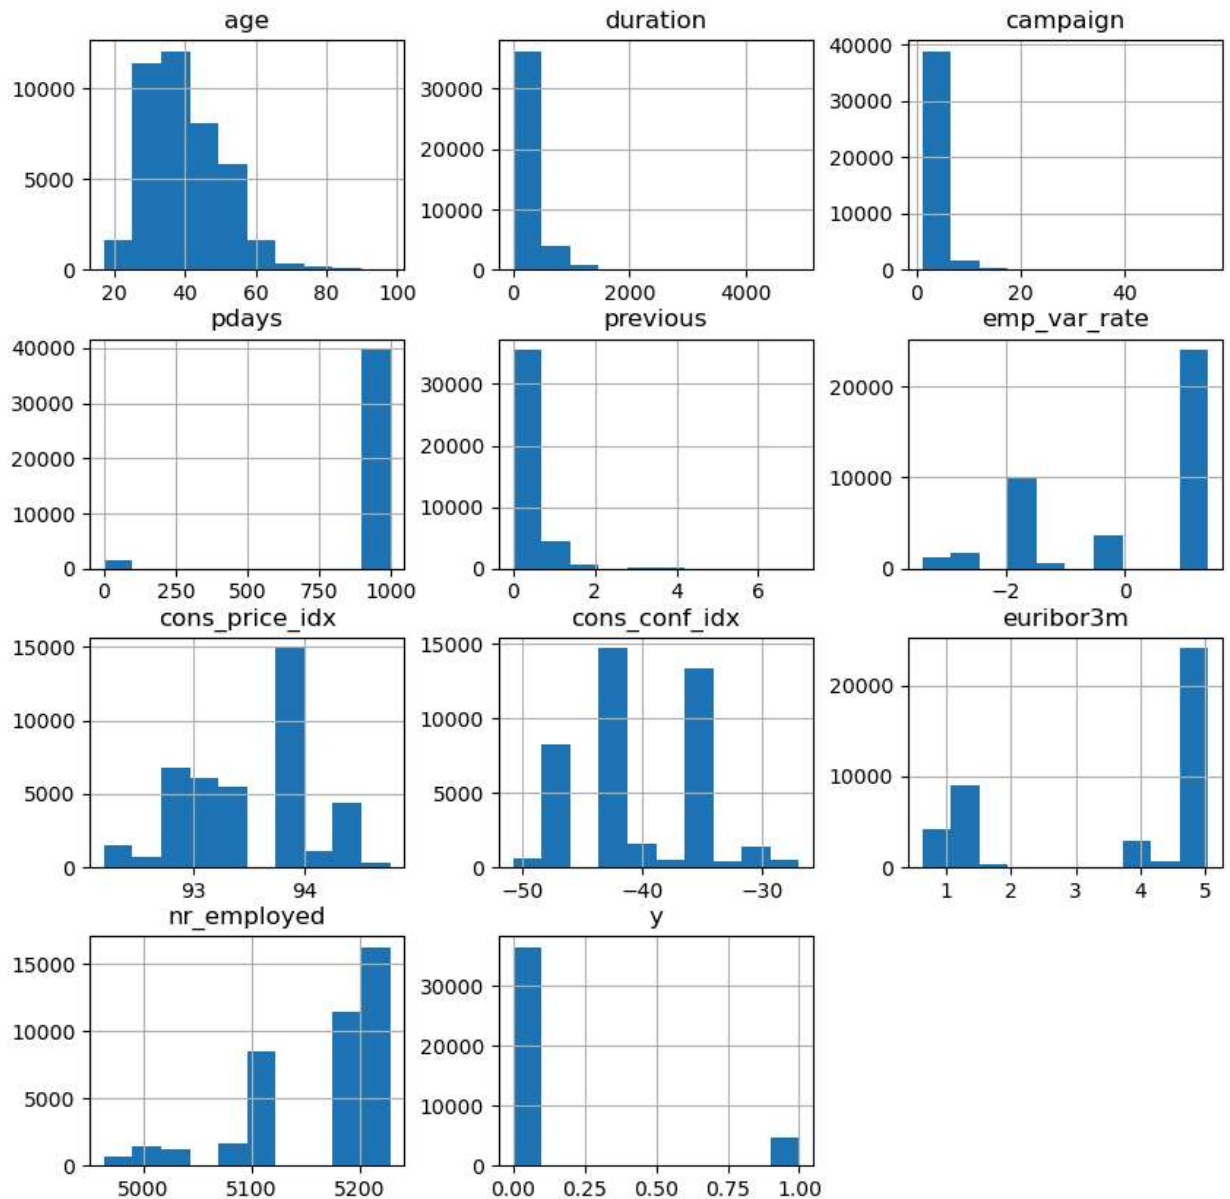
```
In [28]: # Check for missing values
         df.isnull().sum()
```

```
Out[28]: age                0
         job                0
         marital            0
         education          0
         default            0
         housing            0
         loan               0
         contact            0
         month              0
         day_of_week        0
         duration           0
         campaign           0
         pdays             0
         previous           0
         poutcome           0
         emp_var_rate       0
         cons_price_idx     0
         cons_conf_idx      0
         euribor3m          0
         nr_employed        0
         y                  0
         dtype: int64
```

```
In [29]: print(df.duplicated().sum())

12
```

```
In [30]: import matplotlib.pyplot as plt
         # Histograms for numerical features
         df3.hist(figsize=(10, 10))
         plt.show()
```



```
In [31]: # Split the data into features (X) and target (y)
X = df.drop('loan', axis=1)
y = df['loan']
```

```
In [32]: # Generate a synthetic dataset
X, y = make_classification(n_samples=1000, n_features=10, n_informative=5, n_redundant=5,
                           random_state=42)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Random Forest

```
In [34]: # Build a random forest classifier with 100 trees
rd = RandomForestClassifier(n_estimators=100, random_state=1)

# Train the classifier on the training data
rd.fit(X_train, y_train)
```

Out[34]:

▼ RandomForestClassifier

RandomForestClassifier(random_state=1)

```
In [35]: # Make predictions on the test data
y_pred = rd.predict(X_test)

# Evaluate the performance of the classifier
accuracy = rd.score(X_test, y_test)
print("Accuracy: %.2f" % accuracy)
```

Accuracy: 0.95

```
In [36]: # Calculate R² score
r2 = r2_score(y_test, y_pred)
print(f"R² score: {r2:.2f}")

# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Absolute Error: {mae:.2f}")
```

R² score: 0.80
Mean Absolute Error: 0.05

```
In [37]: report = classification_report(y_test, y_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	141
1	0.99	0.91	0.95	159
accuracy			0.95	300
macro avg	0.95	0.95	0.95	300
weighted avg	0.95	0.95	0.95	300

Apply Hyper parameter tuning

```
In [39]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
```

```
In [40]: # Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 200, 300],
    'max_depth': [None, 10, 20, 30, 40],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
```

```
In [41]: # Create a Random Forest model
rf = RandomForestClassifier(random_state=42)

# Set up GridSearchCV
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, n_jobs=-1, verbose=1)
```

```
# Fit GridSearchCV
grid_search.fit(X_train, y_train)
```

```
# Get the best parameters
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)
```

Fitting 5 folds for each of 360 candidates, totalling 1800 fits

Best parameters found: {'bootstrap': False, 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}

```
In [42]: # Define the parameter distributions
param_dist = {
    'n_estimators': [50, 100, 200, 300],
    'max_depth': [None, 10, 20, 30, 40],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

# Set up RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=rf, param_distributions=param_dist, n_iter=100)

# Fit RandomizedSearchCV
random_search.fit(X_train, y_train)

# Get the best parameters
best_params = random_search.best_params_
print("Best parameters found: ", best_params)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

Best parameters found: {'n_estimators': 50, 'min_samples_split': 2, 'min_samples_leaf': 2, 'max_depth': 30, 'bootstrap': False}

```
In [43]: # Retrieve the best model from RandomizedSearchCV
best_rf = random_search.best_estimator_

# Make predictions
y_pred = best_rf.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Test set accuracy: ", accuracy)
```

Test set accuracy: 0.9666666666666667

Adaboost Algorithm

```
In [45]: # Train the AdaBoost classifier
ada = AdaBoostClassifier(n_estimators=50, learning_rate=1.0, random_state=42)
ada.fit(X_train, y_train)

# Make predictions on the test set
y_pred = ada.predict(X_test)

# Calculate accuracy of the classifier
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the AdaBoost classifier: {accuracy * 100:.2f}")
```

Accuracy of the AdaBoost classifier: 91.33

```
In [46]: # Calculate R2 score
r2 = r2_score(y_test, y_pred)
print(f"R2 score: {r2:.2f}")

# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Absolute Error: {mae:.2f}")
```

R² score: 0.65

Mean Absolute Error: 0.09

```
In [47]: report = classification_report(y_test,y_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.86	0.97	0.91	141
1	0.97	0.86	0.91	159
accuracy			0.91	300
macro avg	0.92	0.92	0.91	300
weighted avg	0.92	0.91	0.91	300

Hyper Paramter Tunning

```
In [49]: # Define the parameter distributions for hyperparameter tuning
param_dist = {
    'n_estimators': [50, 100, 200, 300, 500], # Number of estimators
    'learning_rate': [0.001, 0.01, 0.1, 0.5, 1.0, 1.5], # Learning rate
    'algorithm': ['SAMME', 'SAMME.R'] # Algorithm type
}

# Set up RandomizedSearchCV for hyperparameter tuning
random_search = RandomizedSearchCV(
    estimator=ada,
    param_distributions=param_dist,
    n_iter=20, # Number of parameter settings that are sampled
    cv=5, # Number of folds for cross-validation
    n_jobs=-1, # Use all available cores for parallelization
    verbose=2, # Verbosity level
    random_state=42
)

# Fit RandomizedSearchCV on the training data
random_search.fit(X_train, y_train)

# Get the best parameters from RandomizedSearchCV
best_params = random_search.best_params_
print("Best parameters found: ", best_params)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

Best parameters found: {'n_estimators': 200, 'learning_rate': 0.5, 'algorithm': 'SAMME'}

```
In [50]: # Retrieve the best AdaBoost model from RandomizedSearchCV
best_ada = random_search.best_estimator_
```

```
# Make predictions using the best AdaBoost model
y_pred = best_ada.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the tuned AdaBoost classifier: {accuracy * 100:.2f}")
```

Accuracy of the tuned AdaBoost classifier: 91.00